

**FEATURE DRIFT DETECTION VIA ADVERSARIAL
VALIDATION**

NUSRAT BEGUM

**A THESIS SUBMITTED IN PARTIAL FULFILLMENT
OF THE REQUIREMENTS FOR THE DEGREE OF
MASTER OF SCIENCE (COMPUTER SCIENCE)
FACULTY OF GRADUATE STUDIES
MAHIDOL UNIVERSITY
2026**

COPYRIGHT OF MAHIDOL UNIVERSITY

Thesis
entitled
**FEATURE DRIFT DETECTION VIA ADVERSARIAL
VALIDATION**

.....
Ms. Nusrat Begum
Candidate

.....
Asst. Prof. Thanapon Noraset,
Ph.D. (Computer Science)
Major advisor

.....
Assoc. Prof. Suppawong Tuarob,
Ph.D. (Computer Science)
Co-advisor

.....
Assistant Professor Dr. Warin
Krityakiarana
Deputy Dean for Planning and Quality
Assurance
Acting as Dean of the Faculty of Graduate
Studies, Mahidol University
Dean
Faculty of Graduate Studies
Mahidol University

.....
Assoc. Prof. Boonsit Yimwadsana,
Ph.D. (Electrical Engineering)
Program Director
Master of Science Programme
in Computer Science
Faculty of Information and
Communication Technology
Mahidol University

Thesis
entitled
**FEATURE DRIFT DETECTION VIA ADVERSARIAL
VALIDATION**

was submitted to the Faculty of Graduate Studies, Mahidol University
for the degree of Master of Science (Computer Science)

on
February 2026

.....
Ms. Nusrat Begum
Candidate

.....
Assistant Professor Dr. Akara Supratak,
Ph.D. (Computer Science)
Chair

.....
Asst. Prof. Thanapon Noraset,
Ph.D. (Computer Science)
Member

.....
Assoc. Prof. Suppawong Tuarob,
Ph.D. (Computer Science)
Member

.....
Assistant Professor Dr. Warin
Krityakiarana
Deputy Dean for Planning and Quality
Assurance
Acting as Dean of the Faculty of Graduate
Studies, Mahidol University
Dean
Faculty of Graduate Studies
Mahidol University

.....
Dr. Pattanasak Mongkolwat,
Ph.D. (Computer Science)
Dean
Faculty of Information and
Communication Technology
Mahidol University

ACKNOWLEDGEMENTS

Nusrat Begum

FEATURE DRIFT DETECTION VIA ADVERSARIAL VALIDATION

NUSRAT BEGUM 6638009 ITCS/M

M.Sc. (COMPUTER SCIENCE)

THESIS ADVISORY COMMITTEE: THANAPON NORASET, Ph.D., SUPPAWONG TUAROB, Ph.D.

ABSTRACT

Machine learning models deployed in production environments frequently experience distribution drift—changes in input feature distributions over time that can degrade model performance. While discriminative methods like D3 have proven that adversarial validation can accurately detect drift, they suffer from the “black box” problem—providing no insight into the source of the drift. In complex MLOps pipelines, knowing *that* drift occurred is insufficient; engineers must understand *what* drifted to take appropriate corrective action.

This thesis proposes **Explainable Adversarial Drift Detection (EADD)**, a novel framework that extends adversarial validation with SHAP-based root cause analysis. Unlike univariate methods (e.g., KSWIN) that flood users with alerts for every feature, EADD employs a multivariate permutation test to robustly flag significant shifts, followed by an automated diagnostic step that ranks features by their contribution to the drift. This transforms the drift detector from a simple alarm into an intelligent diagnostic tool.

The main contributions are: (1) a streaming adversarial validation framework with reservoir sampling for representative reference windows and statistical rigor through permutation testing; (2) SHAP-based root cause analysis that identifies which specific features drive detected drift; (3) automated drift prescriptions providing actionable remediation recommendations based on drift patterns (univariate, multivariate, or feature-subset drift); (4) comprehensive benchmarking against state-of-the-art methods (D3, ADWIN, KSWIN, MMD) on the Lukats et al. (2025) real-world data streams; and (5) practical implementation guidelines for MLOps deployment including threshold calibration and computational efficiency considerations.

This work establishes EADD as a principled, explainable drift detection methodology that advances beyond detection accuracy to actionable diagnostics, addressing the critical explainability gap identified in recent benchmarks.

IMPLICATION OF THE THESIS

EADD provides both high detection accuracy and actionable explanations, enabling MLOps teams to respond appropriately to different drift scenarios rather than treating all drift signals as equivalent alarms requiring full model retraining.

KEY WORDS : DRIFT DETECTION / ADVERSARIAL VALIDATION /
COVARIATE DRIFT / MLOPS / DISTRIBUTION SHIFT /
STREAMING DATA

38 pages

CONTENTS

	Page
ACKNOWLEDGEMENTS	iii
ABSTRACT	iv
LIST OF TABLES	viii
LIST OF FIGURES	ix
CHAPTER I INTRODUCTION	1
1.1 Motivation	1
1.2 Objectives	2
1.3 Expected Contributions	3
CHAPTER II LITERATURE REVIEW	5
2.1 Overview	5
2.2 Types of Drift	6
2.2.1 Drift by Nature of Change	6
2.2.2 Drift by Temporal Pattern	7
2.3 Classification of Drift Detection Methods	7
2.3.1 Supervised Drift Detection Methods	8
2.3.2 Unsupervised Drift Detection Methods	9
2.3.3 Hybrid and Ensemble Methods	13
2.4 Limitations of Existing Methods	14
2.5 Related Work: Adversarial Validation	15
2.5.1 Concept and Origins	15
2.5.2 Industrial Applications	15
2.6 Gap in the Literature	16
CHAPTER III METHODOLOGY	18
3.1 Overview: Explainable Adversarial Drift Detection (EADD)	18
3.2 Problem Formulation	19
3.3 EADD Framework Components	20

CONTENTS (cont.)

	Page
3.3.1 Step 1: Adaptive Reference Windowing	21
3.3.2 Step 2: Adversarial Validation (The Detector)	21
3.3.3 Step 3: Statistical Significance via Permutation Test	22
3.3.4 Step 4: Feature Drift Explainability via SHAP (Novel Contribution)	23
3.3.5 Automated Drift Prescription (Novel Contribution)	24
3.3.6 Model Adaptation Strategies	25
3.4 Datasets	26
3.5 Baseline Drift Detectors	27
3.6 Evaluation Metrics	27
3.7 Experimental Design	28
3.8 Practical Considerations	29
CHAPTER IV RESULTS	30
CHAPTER V DISCUSSION	31
CHAPTER VI CONCLUSION	32
BIOGRAPHY	38

LIST OF TABLES

Table	Page
3.1 Alignment of EADD Components with Research Requirements	20

LIST OF FIGURES

Figure	Page
2.1 Comparison of no drift versus drift scenarios. Left: When distributions overlap, the adversarial classifier cannot distinguish data sources ($AUC \approx 0.5$). Right: When drift occurs, distributions shift and the classifier can distinguish sources ($AUC \gg 0.5$).	5
3.1 Overview of Explainable Adversarial Drift Detection (EADD) methodology. The framework extends adversarial validation with SHAP-based root cause analysis and automated prescriptions.	18

CHAPTER I

INTRODUCTION

1.1 Motivation

Machine learning models deployed in real-world environments frequently encounter evolving data distributions over time. According to Lu et al. (2018) and Gama et al. (2014), these changes—collectively referred to as distribution shift or drift—can manifest as changes in input features (covariate drift) or in the relationship between inputs and labels (concept drift), causing model performance to deteriorate if left unmanaged. Lu et al. (2018) and Widmer and Kubat (1996) emphasize that detecting such changes early is essential for maintaining reliable model performance, reducing operational risk, and avoiding accumulating technical debt in production systems (Bayram et al., 2022; Sculley et al., 2015).

Conventional model monitoring strategies react to drift only after observable performance degradation, which by design means that the system has already suffered from poor predictions on new data (Widmer and Kubat, 1996). In production contexts where ground-truth labels may be delayed, unavailable, or costly to obtain, this post-hoc reactionary approach is often inadequate (Lu et al., 2018; Souza et al., 2020a). According to Lu et al. (2018) and Bayram et al. (2022), early drift detection that anticipates distribution changes before significant performance loss is therefore a practical priority in dynamic machine learning systems.

Current drift detection methods are varied. Traditional statistical methods employ hypothesis tests—such as the Kolmogorov–Smirnov test, chi-squared test, or Kullback–Leibler divergence—to detect distribution changes by comparing incoming and reference data (Rabanser et al., 2019; Raab et al., 2020). Bifet and Gavaldà (2007) developed adaptive windowing methods like ADWIN that dynamically adjust how much historical data is retained for comparison. Supervised error-based detectors like DDM and EDDM monitor model performance metrics over time, but require access to ground-truth

labels (Gama et al., 2004; Baena-García et al., 2006). Prediction-based detectors track model outputs for changes in confidence, entropy, or prediction distributions without necessarily requiring labels (Sethi and Kantardzic, 2017). Despite this diversity, several challenges remain: statistical tests can be sensitive to noise and may struggle in high-dimensional settings; supervised methods depend on label availability; prediction-based methods may be confounded by model calibration issues; and hybrid approaches often require engineered context variables that may not generalize across domains (Lu et al., 2018; Webb et al., 2016).

To address these gaps, we propose recasting feature drift detection as a discriminative learning problem: treating the comparison between historical and newly arriving data as a binary classification task. If an auxiliary classifier can reliably separate data from the two time periods based solely on features, it indicates a meaningful distributional change. This approach, known as adversarial validation, has been used in data science to detect and quantify dataset shift Pan et al. (2020); Qian and Rao (2021) but has not been systematically developed into a general feature drift detector for streaming or window-based data monitoring.

According to Lopez-Paz and Oquab (2017), this approach offers several key advantages for drift detection: it is label-independent (does not require ground-truth labels for new data), model-agnostic (can be applied regardless of the primary prediction model), makes no assumptions about which features or contexts are important, and can detect complex high-dimensional, non-linear distributional shifts that may evade univariate or marginal statistical tests.

1.2 Objectives

The primary objective of this thesis is to develop **Explainable Adversarial Drift Detection (EADD)**, a novel framework that extends adversarial validation with root cause analysis and actionable prescriptions. Specifically, this work aims to:

1. Formalize feature drift detection as a discriminative two-sample learning problem, extending beyond binary detection to provide **feature-level**

attribution of drift sources using SHAP-based interpretability methods (Lundberg and Lee, 2017).

2. Develop streaming and sliding window procedures with **reservoir sampling** for representative reference windows, and **permutation testing** for statistically principled threshold calibration.
3. Systematically benchmark EADD against state-of-the-art unsupervised drift detectors—including D3 Gözüaık et al. (2019), ADWIN Bifet and Gavaldà (2007), KSWIN Raab et al. (2020), and MMD Gretton et al. (2012)—on the real-world datasets from the Lukats et al. (2025) benchmark Lukats et al. (2025).
4. Design and validate an **automated prescription system** that provides actionable remediation recommendations based on drift diagnosis (univariate drift, multivariate concept shift, or feature-subset drift).
5. Evaluate practical applicability in MLOps settings, including computational efficiency, threshold sensitivity, and integration with model retraining strategies.

1.3 Expected Contributions

This thesis seeks to contribute:

1. **EADD Framework:** A formal explainable adversarial validation framework that combines discriminative drift detection with SHAP-based root cause analysis, transforming drift detection from a binary alarm into an intelligent diagnostic tool.
2. **Root Cause Analysis:** SHAP-based feature importance extraction that identifies which specific features drive detected drift, addressing the “black box” limitation of existing methods like D3 (Gözüaık et al., 2019).

3. **Automated Prescriptions:** A novel classification of drift patterns (univariate, multivariate, feature-subset) with corresponding actionable remediation recommendations for MLOps teams.
4. **Benchmark Evaluation:** Comprehensive comparison with state-of-the-art unsupervised detectors on real-world data streams, demonstrating both detection accuracy and the value of explainability.
5. **Practical Guidelines:** Deployment recommendations for MLOps pipelines, including permutation-based threshold calibration, reservoir sampling for reference windows, and computational efficiency considerations.

CHAPTER II

LITERATURE REVIEW




Figure 2.1: Comparison of no drift versus drift scenarios. Left: When distributions overlap, the adversarial classifier cannot distinguish data sources ($AUC \approx 0.5$). Right: When drift occurs, distributions shift and the classifier can distinguish sources ($AUC > 0.5$).

2.1 Overview

Drift detection addresses the challenge that arises when the statistical characteristics of data evolve over time in a manner that impacts the validity of a machine learning model (Lu et al., 2018; Gama et al., 2014). According to Bayram et al. (2022), when users, environments, data sources, or the relationships between labels and features change, model accuracy can diminish. Gama et al. (2014) and Sculley et al. (2015) emphasize that drift detection identifies this modification early, allowing practitioners to adapt, retrain, or update the model accordingly. A robust understanding of drift types and detection methods is essential for designing effective monitoring systems.

2.2 Types of Drift

Drift can be categorized along two main dimensions: by the nature of the change (what component of the data-generating process is affected) and by the temporal pattern of the change (how the drift manifests over time) (Lu et al., 2018; Gama et al., 2014).

2.2.1 Drift by Nature of Change

Based on which component of the joint distribution $P(X, y)$ changes, the literature (Lu et al., 2018; Gama et al., 2014; Tsymbal, 2004) classifies drift into the following types:

Concept Drift (Real Drift): The conditional distribution $P(y|X)$ changes over time while the feature distribution $P(X)$ may or may not change (Lu et al., 2018; Žliobaitė, 2010). This represents a change in the underlying relationship between inputs and outputs—the “concept” the model has learned becomes invalid. For example, the factors that determine creditworthiness may evolve due to economic shifts, even if the applicant characteristics (features) remain similar (Gama et al., 2014). Concept drift directly degrades model performance because the learned decision boundaries no longer reflect the true input-output relationship (Lu et al., 2018).

Covariate Drift (Feature Drift / Virtual Drift): The marginal distribution of input features $P(X)$ changes over time while the conditional relationship $P(y|X)$ remains stable (Tsymbal, 2004; Shimodaira, 2000). The ground-truth concept does not change, but the model may encounter inputs outside its training distribution. For example, a fraud detection model trained on transaction patterns from one demographic may receive data from a different user population, even though the definition of fraud remains unchanged (Tsymbal, 2004). Covariate drift can degrade model performance because the model is exposed to feature regions where it may not generalize well (Shimodaira, 2000).

Prior Probability Drift (Label Shift): The marginal distribution of labels $P(y)$ changes over time while $P(X|y)$ remains stable (Tsymbal, 2004). This occurs when the prevalence of different classes changes—for instance, a seasonal increase in a particular disease affecting the class balance in medical diagnosis. This is sometimes called class imbalance shift (Lu et al., 2018).

Note on Terminology: The literature uses varied terminology (Lu et al., 2018). The term “concept drift” is sometimes used as an umbrella term for any distribution change affecting model performance (Gama et al., 2014), while some reserve it specifically for changes in $P(y|X)$ (Tsymbal, 2004). This thesis focuses primarily on covariate (feature) drift detection, as it enables early warning without requiring labels.

2.2.2 Drift by Temporal Pattern

Based on how drift manifests over time, the literature (Lu et al., 2018; Gama et al., 2014; Žliobaitė, 2010) categorizes drift as follows:

- **Abrupt (Sudden) Drift:** A rapid, discrete change where the data distribution shifts quickly from one state to another within a short time period (Gama et al., 2014).
- **Gradual Drift:** A new distribution progressively replacing the old one over an extended duration, with samples from both distributions interleaved during the transition period (Lu et al., 2018).
- **Incremental Drift:** Small but continuous changes accumulating over time, eventually resulting in a significant distributional shift (Žliobaitė, 2010).
- **Recurring (Seasonal) Drift:** Previously observed distributions reappearing periodically due to cyclical patterns, such as seasonal effects in retail or recurring user behavior patterns (Lu et al., 2018; Gama et al., 2014).

Understanding drift patterns is important for detector design: abrupt drift requires rapid response, gradual drift requires sustained sensitivity, and recurring drift benefits from mechanisms that can recognize previously seen distributions (Gama et al., 2014).

2.3 Classification of Drift Detection Methods

Based on comprehensive surveys of drift detection literature (Lu et al., 2018; Gama et al., 2014; Tsymbal, 2004; Hinder et al., 2024), drift detection methods can be

classified by multiple dimensions. According to Lu et al. (2018), the most fundamental distinction is based on label availability: whether the method requires access to ground-truth labels (supervised) or operates without them (unsupervised).

2.3.1 Supervised Drift Detection Methods

According to Gama et al. (2014) and Gama et al. (2004), supervised drift detection methods require access to ground-truth labels and primarily monitor model performance or prediction errors to detect distributional changes. Gama et al. (2004) and Baena-García et al. (2006) explain that these methods operate under the assumption that unexpected changes in error characteristics—such as increases in error rate, loss magnitude, variance, or cumulative deviation—are indicative of underlying drift.

Key supervised drift detection algorithms include:

Drift Detection Method (DDM): According to Gama et al. (2004), DDM monitors the online error rate and its standard deviation. When these statistics exceed predefined thresholds derived from binomial distribution bounds, DDM signals a warning or drift state. Gama et al. (2004) note that DDM is effective for abrupt drift but may be slow to detect gradual changes.

Early Drift Detection Method (EDDM): EDDM extends DDM by tracking the distance between consecutive classification errors rather than just the error rate (Baena-García et al., 2006). This modification improves sensitivity to gradual drift by detecting when errors cluster more closely together.

Page–Hinkley Test: This sequential analysis technique uses cumulative sums to detect changes in the mean of a monitored signal (Page, 1954). When cumulative deviation exceeds a threshold, drift is indicated.

Hoeffding Drift Detection Methods (HDDM-A, HDDM-W): These methods use Hoeffding’s inequality to provide statistically grounded drift decisions with theoretical guarantees (Frias-Blanco et al., 2015). HDDM-A detects abrupt changes using moving averages, while HDDM-W is designed for gradual drift using window-based statistics.

Limitations: Supervised methods fundamentally require labeled data, which may be delayed, expensive, or unavailable in many production scenarios (Lu et al.,

2018; Sethi and Kantardzic, 2017). They also react to drift after it has already caused performance degradation, rather than detecting distributional change proactively.

2.3.2 Unsupervised Drift Detection Methods

Unsupervised drift detection methods do not require ground-truth labels and instead monitor changes in data distribution, model outputs, or internal representations (Tsymbal, 2004; Hinder et al., 2024). These methods enable earlier detection of distributional shifts—before performance degradation becomes observable—making them particularly valuable in production environments with label delays (Sethi and Kantardzic, 2017).

2.3.2.1 Statistical Distribution-Based Methods

Statistical test-based methods detect drift by directly examining changes in data distributions through hypothesis testing or distance-based measures (Rabanser et al., 2019; Raab et al., 2020). A common strategy compares a reference data window against a current data window to test whether they originate from the same distribution (Bifet and Gavalda, 2007).

Key statistical drift detection methods include:

Adaptive Windowing (ADWIN): ADWIN dynamically maintains a sliding window and automatically shrinks it when statistically significant distributional change is detected (Bifet and Gavalda, 2007). It provides theoretical guarantees on false positive and false negative rates and adapts its window size based on observed change.

KSWIN: This method applies the Kolmogorov–Smirnov two-sample test over sliding windows to detect changes in one-dimensional feature distributions (Raab et al., 2020). KSWIN is non-parametric and makes minimal assumptions about the data distribution.

Chi-Squared Test-Based Detectors: These use the chi-squared test to compare categorical feature distributions or discretized continuous features between time periods (Lu et al., 2018).

Kullback–Leibler (KL) Divergence and Other Divergence Measures: Methods that quantify the “distance” between probability distributions, including KL divergence, Jensen–Shannon divergence, and Hellinger distance, can all be used to

measure distributional shift (Dasu et al., 2006).

Maximum Mean Discrepancy (MMD): MMD is a kernel-based two-sample test that compares distributions by mapping them into a reproducing kernel Hilbert space, capable of capturing complex multivariate distributional differences (Gretton et al., 2012).

Strengths and Limitations: Statistical tests are interpretable, theoretically grounded, and label-free (Rabanser et al., 2019). However, they can be sensitive to noise (causing false alarms), may struggle in high-dimensional spaces due to the curse of dimensionality, and typically test marginal distributions rather than joint distributions (Lu et al., 2018; Webb et al., 2016).

2.3.2.2 Model-Based Discriminative Methods

According to Pan et al. (2020), Lopez-Paz and Oquab (2017), and Sethi and Kantardzic (2017), discriminative drift detection methods train an auxiliary model to distinguish between reference and current data distributions. Lopez-Paz and Oquab (2017) explain that the core idea is to frame drift detection as a classification problem: if a classifier can accurately predict whether each instance originates from the reference or current distribution, the distributions must be distinguishably different.

Key approaches include:

Domain Classifier / Adversarial Validation: According to Pan et al. (2020) and Qian and Rao (2021), this approach trains a binary classifier (e.g., logistic regression, random forest, gradient boosting) to discriminate samples from reference vs. current windows. Pan et al. (2020) note that high discriminative performance (measured by accuracy, AUC, or similar metrics) indicates significant distributional shift.

Discriminative Drift Detector (D3): Gözüaık et al. (2019) developed this unsupervised method that trains a classifier (logistic regression) to distinguish between old and new data windows, using AUC as the drift signal. D3 has demonstrated strong performance in recent benchmarks (Lukats et al., 2025) and represents the state-of-the-art in discriminative drift detection.

Margin Density Drift Detection (MD3): Sethi and Kantardzic (2017) proposed an alternative approach that monitors margin density—the proportion of samples near the decision boundary—to detect drift without requiring true labels.

Margin-Based and Density-Based Discriminators: According to Ditzler

and Polikar (2011), these methods monitor changes in classifier margin densities or decision boundary proximity to detect distributional shifts.

Strengths: Discriminative methods can capture high-dimensional, non-linear distributional shifts that marginal univariate tests may miss (Lopez-Paz and Oquab, 2017). They require no labels for the drift detection task itself (only pseudo-labels indicating data source) and naturally aggregate information across all features (Pan et al., 2020).

Limitations: These methods require training a classifier at each monitoring step, which has computational cost; performance depends on classifier choice and hyper-parameters; and interpretation of drift signals may be less intuitive than statistical test p-values (Sethi and Kantardzic, 2017).

2.3.2.3 Representation-Based Methods

Representation-based methods detect drift by monitoring changes in learned representations or embeddings rather than raw feature distributions (Rabanser et al., 2019; Liu et al., 2020):

Embedding Distribution Shift: Computing statistical distances (cosine similarity, Wasserstein distance, MMD) between reference and current embedding distributions derived from neural network intermediate layers (Rabanser et al., 2019).

Deep Learning Representation Monitoring: This approach tracks stability of learned representations across time, leveraging insights from domain adaptation and transfer learning literature (Long et al., 2016).

These methods are particularly relevant for deep learning systems where raw feature distributions may not reflect semantically meaningful changes (Liu et al., 2020).

2.3.2.4 Prediction-Based / Output-Based Methods

Prediction-based methods detect drift by monitoring changes in model outputs over time without requiring evaluation of prediction correctness (Sethi and Kantardzic, 2017; Gözüaık et al., 2019):

Prediction Distribution Shift: Monitoring changes in the distribution of model predictions (class probabilities, regression outputs) over time (Gözüaık et al., 2019).

Uncertainty / Entropy Monitoring: Tracking prediction entropy or uncertainty scores under the assumption that rising uncertainty reflects exposure to out-of-

distribution data (Guo et al., 2017).

Confidence Monitoring: Tracking changes in model confidence levels and softmax output distributions (Guo et al., 2017).

Limitations: These methods depend on model calibration. If the model is poorly calibrated—which is common for modern deep learning models—changes in confidence may not reliably indicate drift (Guo et al., 2017; Ovadia et al., 2019).

2.3.2.5 State-of-the-Art Unsupervised Detectors (2019–2025)

Recent benchmarking by Lukats et al. (2025) provides a systematic evaluation of fully unsupervised drift detectors on real-world data streams. Their study tested numerous algorithms and identified seven that function reliably on messy, real-world data. This survey represents a critical “reality check” for the field, demonstrating that detecting change without ground-truth labels is possible but fragile.

D3 (Discriminative Drift Detector): According to Gözüaık et al. (2019), D3 trains a logistic regression classifier to distinguish between old and new data windows, using AUC as the drift signal. Lukats et al. (2025) found D3 to be the overall winner—the most robust and accurate across different data types. However, D3 remains a “black box”: it signals *that* drift occurred but provides no insight into *which* features drifted or *why*.

CSDDM (Clustered Statistical Drift Detection Method): According to Hsieh et al. (2021), this method uses PCA for dimensionality reduction followed by K-Means clustering, then applies statistical tests to detect cluster shifts. Lukats et al. (2025) report it performed exceptionally well on the “lift-per-drift” metric, catching drifts with minimal false alarms.

SPLL (Semi-Parametric Log-Likelihood): Kuncheva (2013) developed this approach that fits probability density models to reference data and detects drift when new observations have anomalously low likelihood under the learned distribution. It is effective for sudden drift patterns.

IBDD (Image-Based Drift Detector): According to Souza et al. (2020b), this creative approach converts numerical data streams into images and applies image comparison techniques to detect distributional changes. While highly accurate, Lukats et al. (2025) note it can be biased by its frequent reset strategy.

OCDD (One-Class Drift Detector): Gözüaık and Can (2021) proposed

using one-class SVMs to define “normal” data boundaries and flag drift when incoming data falls outside these boundaries. While sound in principle, D3 generally outperformed OCDD in benchmarks (Lukats et al., 2025).

NN-DVI (Nearest Neighbor Density Variation): According to Liu et al. (2018), this method monitors changes in local data density by tracking nearest-neighbor distances. Drift is signaled when density patterns shift significantly. It performs well for gradual density changes but can be sensitive to noise.

UCDD (Unsupervised Concept Drift Detection): Shang et al. (2020) developed a clustering-based approach that assigns pseudo-labels based on cluster membership and monitors boundary shifts. It serves as a reliable baseline in unsupervised drift detection.

Industrial Drift Detection Libraries: Beyond research algorithms, several industrial libraries provide production-ready drift detection:

- **Alibi Detect:** Python library implementing Kolmogorov-Smirnov tests, MMD, and learned kernel methods for multivariate drift detection (Liu et al., 2020).
- **NannyML:** Confidence-based Performance Estimation (CBPE) for monitoring model performance without ground-truth labels (Bayram et al., 2022).

Gap Identified from Benchmark Analysis: While these libraries and algorithms exist, according to Lukats et al. (2025), they share a common limitation: they typically report only *whether* drift occurred, not *where* or *why*. Statistical libraries often run separate univariate tests for each feature, causing “false alarm storms” when many features are monitored simultaneously. The field lacks methods that provide a unified multivariate check with granular explanations of drift sources.

2.3.3 Hybrid and Ensemble Methods

Hybrid approaches integrate multiple sources of information—both supervised and unsupervised signals—to improve drift detection robustness (Gama et al., 2014; Krawczyk et al., 2017):

Context-Aware Drift Detection: Combining model predictions, user context, operational signals, and metadata to compute context-weighted drift scores (Krawczyk et al., 2017).

Multi-Signal Ensemble Methods: Integrating predictions, embeddings, operational KPIs, and metadata into a unified drift score (Krawczyk et al., 2017).

Drift Detection Ensembles: Combining multiple base detectors using voting, weighting, or meta-learning to reduce both false alarms and missed detections (Krawczyk et al., 2017).

2.4 Limitations of Existing Methods

While the drift detection literature offers diverse approaches, each category has characteristic limitations:

Supervised Methods (DDM, EDDM, HDDM): These require timely labeled data, which is often delayed or unavailable in production (Lu et al., 2018). They detect drift reactively—only after performance degradation has already occurred (Sethi and Kantardzic, 2017).

Statistical Distribution Methods (ADWIN, KSWIN, KL-Divergence): These can be sensitive to noise, leading to false alarms (Rabanser et al., 2019). They may struggle in high-dimensional or correlated feature spaces and typically test marginal distributions, potentially missing joint distribution changes (Lu et al., 2018; Webb et al., 2016).

Prediction/Uncertainty-Based Methods: These rely on model calibration. Modern ML models are often poorly calibrated, especially under distribution shift, causing these signals to be unreliable (Guo et al., 2017; Ovadia et al., 2019).

Context-Aware Methods: These require pre-specification of relevant context variables, which may be incomplete, noisy, or themselves drift over time (Krawczyk et al., 2017). Misspecified context can cause missed detections or spurious alarms.

Fundamental Gap: Critically, most existing methods do not directly test whether new data is statistically distinguishable from historical data in the original feature

space (Lopez-Paz and Oquab, 2017). Discriminative approaches offer a natural, model-agnostic signal of drift that does not depend on predefined contexts, labels, or calibrated uncertainties—yet this perspective has received limited systematic attention in the drift detection literature (Pan et al., 2020).

This thesis aims to bridge this gap by establishing adversarial validation as a viable, general-purpose drift detection technique through rigorous empirical evaluation and practical methodology development.

2.5 Related Work: Adversarial Validation

2.5.1 Concept and Origins

According to Pan et al. (2020), adversarial validation was originally proposed as a practical technique for diagnosing dataset shift between training and evaluation data. Pan et al. (2020) and Qian and Rao (2021) explain that the approach reframes the comparison of two datasets as a binary classification problem: train a classifier to predict whether each instance originates from the reference distribution (e.g., training data) or the target distribution (e.g., test or new data). Lopez-Paz and Oquab (2017) note that high discriminative performance indicates statistical separability—hence significant distributional shift—while near-random performance suggests the datasets are largely indistinguishable.

According to Pan et al. (2020), the term “adversarial” reflects the insight that if an adversary (the classifier) cannot distinguish the data sources, the distributions must be similar. Lopez-Paz and Oquab (2017) note this is related to the discriminator component in Generative Adversarial Networks (GANs), though applied here for diagnostic rather than generative purposes.

2.5.2 Industrial Applications

Uber’s MaLTA System: Pan et al. (2020) presented adversarial validation at Uber for their MaLTA (Machine Learning Targeting Automation) user targeting system. They trained a classifier to distinguish historical training data from newly observed targeting data. When the adversarial classifier exhibited strong separability, it signaled

train-deployment mismatch, prompting feature filtering and retraining adjustments before deployment. Their results demonstrated improved robustness compared to naive retraining strategies that react only after observable performance degradation.

Credit Scoring: Qian and Rao (2021) applied adversarial validation to manage dataset shift in credit scoring. They trained an adversarial classifier to distinguish historical credit applications from recent ones, using the classifier’s predictions to filter training samples—keeping only those most similar to the current deployment distribution for model training.

Image Classification and Operations: Additional applications include using adversarial validation to quantify dataset complexity in image classification Renza and Moya-Albor (2024) and to detect when historical data is no longer representative in manufacturing/operations settings Senoner et al. (2023).

2.6 Gap in the Literature

While discriminative methods like D3 have proven that adversarial validation can accurately detect drift (Gözüaçık et al., 2019; Lukats et al., 2025), significant gaps remain that limit their practical utility in production MLOps environments:

1. **The “Black Box” Problem:** According to Lukats et al. (2025), existing discriminative detectors like D3 signal *that* drift occurred but provide no insight into *which* features drifted or *why*. In complex ML systems, knowing that drift happened is insufficient—engineers need to understand *what drifted* to take appropriate corrective action.
2. **No Root Cause Analysis:** Current methods function like a car’s “Check Engine” light—they indicate a problem exists but not which component failed (Lukats et al., 2025). Production systems require diagnostic capabilities that pinpoint drift sources.
3. **False Alarm Storm Problem:** Industrial libraries often run separate univariate tests for each feature, causing multiple simultaneous alerts

when features are correlated (Rabanser et al., 2019). This overwhelms operators rather than guiding them toward the actual drift source.

4. **Threshold Calibration Challenges:** According to Lukats et al. (2025), determining appropriate detection thresholds remains a major unsolved problem. Existing methods either use arbitrary thresholds or require labeled validation data that may not be available.
5. **No Actionable Prescriptions:** Even when drift is detected, current methods provide no guidance on appropriate response strategies—whether to retrain fully, adjust specific features, or apply targeted corrections (Bayram et al., 2022).

According to recent benchmarks (Lukats et al., 2025), D3 represents the state-of-the-art in detection accuracy. However, its lack of explainability severely limits practical deployment. This thesis proposes **Explainable Adversarial Drift Detection (EADD)**, which extends adversarial validation with SHAP-based root cause analysis (Lundberg and Lee, 2017) to transform drift detection from a simple alarm into an intelligent diagnostic tool.

This thesis aims to bridge this explainability gap by developing a framework that not only detects drift with high accuracy but also provides actionable insights about drift sources and appropriate remediation strategies.

CHAPTER III

METHODOLOGY




Figure 3.1: Overview of Explainable Adversarial Drift Detection (EADD) methodology. The framework extends adversarial validation with SHAP-based root cause analysis and automated prescriptions.

3.1 Overview: Explainable Adversarial Drift Detection (EADD)

This thesis proposes **Explainable Adversarial Drift Detection (EADD)**, a framework designed specifically for **Feature Drift** ($P(X)$ shift). The fundamental insight is that **Adversarial Validation** and **Feature Drift Detection** are mathematically equivalent—both ask: “*Can a model distinguish between old data and new data based solely on features?*”

While existing methods like the Discriminative Drift Detector (D3) (Gözüağık et al., 2019) successfully utilize adversarial validation to flag distributional changes, they function as “black boxes”—providing a drift signal without explaining *which* features drifted. Current drift detection methods function like a car’s “Check Engine” light; EADD is designed as a “Diagnostic Scanner” that identifies exactly which features contributed to the drift, enabling targeted corrective action.

EADD extends the D3 architecture with three key innovations:

1. **Permutation Test:** Statistical validation ensuring detected drift is significant, not noise
2. **SHAP-based Root Cause Analysis:** Granular feature attribution identifying *which* features drifted
3. **Automated Prescriptions:** Actionable remediation recommendations based on drift diagnosis

3.2 Problem Formulation

According to Pan et al. (2020) and Lopez-Paz and Oquab (2017), this research formalizes feature (covariate) drift detection as a problem of discriminating between historical reference data and newly arriving observations in a streaming or windowed setting. Shimodaira (2000) explains that the objective is to detect when the underlying feature distribution $P(X)$ changes substantially over time, providing an early indication of distributional drift before observable degradation in task performance occurs.

Formal Definition: Let $\{x_t\}_{t=1}^T$ denote a stream of feature vectors arriving over time. At each monitoring point t , a drift detector produces a binary decision $d_t \in \{\text{no-drift}, \text{drift}\}$ based on comparison between:

- A reference window W_{ref} containing historical samples presumed to represent the stable/baseline distribution
- A current window W_{cur} containing the most recent observations under test

Extended Objective: Unlike existing detectors that output only a binary drift decision, EADD additionally produces:

- A ranked list of features contributing to the drift
- Feature importance scores quantifying each feature's contribution
- An automated prescription for appropriate remediation action

According to Lu et al. (2018) and Gama et al. (2014), the detection task must balance:

- **Promptness:** Low detection delay (lag between true drift and detection)
- **Reliability:** Low false positive rate (false alarms when no drift exists)
- **Sensitivity:** High true positive rate (detecting drift when it occurs)
- **Efficiency:** Bounded computational and memory requirements suitable for production
- **Explainability:** Clear attribution of drift to specific features (novel contribution)

3.3 EADD Framework Components

The EADD methodology maps directly to the thesis requirements:

Table 3.1: Alignment of EADD Components with Research Requirements

Requirement	EADD Solution
“Based on Adversarial Validation”	Step 2: Train classifier to distinguish W_{ref} vs W_{cur}
“Based on Permutation Test”	Step 3: Shuffle labels $100\times$ to calculate p-value
“Like D3”	Foundation: Same discriminative logic, upgraded classifier (LightGBM)
“Focus on Feature Drift”	Step 4: SHAP identifies which $P(X)$ components shifted

The methodology consists of four integrated steps:

Step 1: Adaptive Reference Windowing — Reservoir sampling for representative baseline

Step 2: Adversarial Validation (Detector) — Train discriminative classifier on W_{ref} vs W_{cur}

Step 3: Permutation Test (Significance) — Statistical validation via label shuffling

Step 4: Feature Drift Explainability (SHAP) — Identify which features drifted (novel contribution)

3.3.1 Step 1: Adaptive Reference Windowing

To detect Feature Drift, we must compare the current data stream against a stable baseline. EADD maintains two windows consistent with established methodologies (Bifet and Gavalda, 2007; Tsymbal, 2004):

Reference Window (W_{ref}): A buffer containing N_{ref} historical samples representing the stable feature distribution $P_{ref}(X)$. Unlike fixed sliding windows that “forget” normal data too quickly, EADD uses **Reservoir Sampling** (Vitter, 1985) to ensure the reference captures the global distribution of features, not just the most recent batches.

Current Window (W_{cur}): A buffer containing the N_{cur} most recent observations representing the current feature distribution $P_{cur}(X)$. As new data arrives, older samples are discarded to maintain responsiveness to recent changes.

Update Mechanism: At each monitoring step:

- New samples enter W_{cur}
- Oldest samples exit W_{cur}
- Reservoir sampling maintains representative W_{ref}
- Upon confirmed drift, W_{ref} may be reset to current distribution

According to Lu et al. (2018), window sizes represent a bias-variance tradeoff: larger windows provide more stable estimates but slower response; smaller windows enable faster detection but higher variance.

3.3.2 Step 2: Adversarial Validation (The Detector)

This step implements the core **Adversarial Validation** logic that forms the foundation of both D3 (Gözüa ık et al., 2019) and EADD. The insight is that Feature Drift detection reduces to a binary classification problem: *can a model tell the difference between old and new data?*

Construct Binary Classification Dataset:

- Label all samples from W_{ref} as class 0 (“old” / reference)
- Label all samples from W_{cur} as class 1 (“new” / current)
- Combined dataset: $X = W_{ref} \cup W_{cur}$, $y = \{0\}^{N_{ref}} \cup \{1\}^{N_{cur}}$

Train Adversarial Classifier: Train a binary classifier f_θ to predict data source from features alone. EADD uses **LightGBM** as the default classifier.

Why LightGBM over D3’s Logistic Regression? D3 (Gözüaık et al., 2019) uses logistic regression, which is linear. LightGBM captures **non-linear feature interactions**, enabling detection of complex Feature Drifts where multiple features shift in correlated ways that linear models miss.

Measure Discriminative Performance: Evaluate classifier performance using stratified cross-validation:

- **AUC-ROC:** Primary metric—Area under the ROC curve
- $AUC \approx 0.5$: Classifier cannot distinguish W_{ref} from $W_{cur} \rightarrow$ **No Feature Drift**
- $AUC \gg 0.5$: Classifier separates windows $\rightarrow P_{ref}(X) \neq P_{cur}(X) \rightarrow$ **Feature Drift detected**

Objective: If the adversarial classifier achieves high AUC, the feature distributions $P(X)$ have diverged—this is the definition of Feature Drift.

3.3.3 Step 3: Statistical Significance via Permutation Test

A raw AUC score (e.g., 0.65) is ambiguous: *Is it drift, or is it noise?* Standard adversarial validation (Pan et al., 2020) uses fixed thresholds (e.g., $AUC > 0.7$), but this is arbitrary. EADD solves this with a **Permutation Test** (Lopez-Paz and Oquab, 2017).

The Permutation Test Procedure:

1. **Generate Null Distribution:** Randomly shuffle the labels (0 and 1) of the combined dataset $B = 100$ times
2. **Retrain:** For each permutation, retrain the adversarial classifier and compute AUC

3. **Build Null:** The 100 permuted AUCs form the “Null Distribution” of AUC scores where *no drift exists*
4. **Compare:** Compare the *actual* adversarial AUC against this Null Distribution

Drift Confirmation Rule:

- Compute empirical p-value: $p = \frac{\#\{AUC_{perm} \geq AUC_{actual}\}}{B}$
- Drift is confirmed **only** if $p < 0.01$ (99th percentile)

Benefit: This prevents false alarms caused by small sample sizes or random fluctuations—a common failure mode in standard adversarial validation. The permutation test provides statistical guarantees that the detected drift is real, not noise.

3.3.4 Step 4: Feature Drift Explainability via SHAP (Novel Contribution)

Once the Permutation Test confirms drift ($p < 0.01$), the critical question becomes: *Which features drifted?* A key limitation of D3 (Gözüaçık et al., 2019) is that it discards the trained classifier after computing AUC. EADD’s key insight is that **this discard is wasteful**—the adversarial classifier has learned exactly *what* makes new data different from old data.

The SHAP Extraction Method: EADD applies **SHAP (SHapley Additive exPlanations)** (Lundberg and Lee, 2017) to the adversarial classifier:

1. Compute SHAP values using TreeSHAP (efficient for LightGBM)
2. Calculate mean absolute SHAP value per feature: $\text{Importance}_i = \frac{1}{n} \sum_{j=1}^n |\text{SHAP}_{ij}|$
3. Normalize to percentages: $\text{Contribution}_i = \frac{\text{Importance}_i}{\sum_k \text{Importance}_k} \times 100\%$
4. Rank features by contribution

Interpretation—The Core Insight:

A feature with a high SHAP value is a feature that **heavily helped the model distinguish “new” data from “old” data**. Therefore: **High SHAP = High Feature Drift**.

Output Format: The system outputs a ranked list of “Drifting Features” with contribution percentages:

- Example: “Drift Detected (AUC=0.85, $p < 0.01$). Primary Driver: **Age** (contribution: 45%), Secondary Driver: **Income** (contribution: 25%), Other: Location (15%), Tenure (10%), Balance (5%)”

This transforms the drift detector from a binary alarm into an **intelligent diagnostic tool** that pinpoints exactly which features require attention.

3.3.5 Automated Drift Prescription (Novel Contribution)

Beyond diagnosis, EADD provides actionable remediation recommendations based on the drift pattern identified:

Scenario A: Univariate Drift (Single Feature Dominant)

Condition: One feature contributes $>80\%$ of total importance.

Prescription: “Univariate drift detected on [Feature]. Recommended actions:

- Investigate data pipeline for [Feature] (sensor failure, data source change)
- Consider feature-specific preprocessing adjustment
- If persistent: Retrain with reduced weight on [Feature] or exclude temporarily”

Scenario B: Multivariate Concept Shift (Distributed Importance)

Condition: No single feature exceeds 30% importance; drift spread across multiple features.

Prescription: “Multivariate concept shift detected. The entire data structure has changed. Recommended actions:

- Full model retraining required immediately
- Consider expanding training data to include recent distribution
- Review deployment assumptions and feature engineering”

Scenario C: Feature Subset Drift

Condition: 2-5 features together contribute $>70\%$ of importance.

Prescription: “Correlated feature drift detected in [Feature Group]. Recommended actions:

- Investigate common data source for affected features
- Consider partial retraining with emphasis on affected feature subspace
- Evaluate feature engineering for the affected domain”

3.3.6 Model Adaptation Strategies

According to Gama et al. (2014), upon drift detection, the primary prediction model can be adapted:

Strategy 1: Drift-Triggered Retraining

- Model is trained at initialization
- Retraining occurs only when drift alarm is raised
- Reference window resets to current data after retraining
- According to Pan et al. (2020), advantage: Conserves resources; exploits early detection for timely updates

Strategy 2: Periodic Retraining (Baseline)

- Model is retrained at fixed intervals regardless of drift signals
- Advantage: Simple; does not depend on detector reliability
- According to Gama et al. (2014), disadvantage: May retrain unnecessarily or too late

Strategy 3: Continuous Update with Drift Weighting

- Online learning with sample weighting based on drift score
- According to Krawczyk et al. (2017), recent samples receive higher weights when drift is detected

3.4 Datasets

According to Souza et al. (2020a), experiments will use both synthetic and real-world benchmark datasets commonly employed in drift detection research:

Synthetic Datasets:

- SEA Concepts (Street and Kim, 2001): Classic drift benchmark with abrupt concept drift
- Rotating Hyperplane: According to Lu et al. (2018), incremental drift through gradual rotation
- STAGGER Concepts: According to Gama et al. (2014), abrupt drift between different rule sets
- Sine/Circle generators: According to Souza et al. (2020a), controlled covariate drift scenarios
- Custom synthetic data: Controlled $P(X)$ shifts with known ground truth

Real-World Benchmark Datasets:

- Electricity (Elec2) (Harries, 1999): Electricity price prediction with natural drift
- Airline: According to Souza et al. (2020a), flight delay prediction with temporal drift
- Forest Covertype: Land cover classification with spatial/temporal variation
- Credit Card Fraud: Transaction data with evolving fraud patterns
- Gas Sensor Array Drift (Vergara et al., 2012): Sensor readings with documented drift

According to Souza et al. (2020a), datasets will be processed in streaming fashion with features used for drift detection and labels reserved for downstream performance evaluation.

3.5 Baseline Drift Detectors

Adversarial validation will be compared against established baselines:

Supervised Baselines:

- DDM (Drift Detection Method) Gama et al. (2004)
- EDDM (Early Drift Detection Method) Baena-García et al. (2006)
- HDDM-A/W (Hoeffding Drift Detection Method) Frias-Blanco et al. (2015)
- Page-Hinkley Test Page (1954)

Unsupervised Statistical Baselines:

- ADWIN (Adaptive Windowing) Bifet and Gavaldà (2007)
- KSWIN (Kolmogorov-Smirnov Windowing) Raab et al. (2020)
- MMD (Maximum Mean Discrepancy) two-sample test Gretton et al. (2012)
- Chi-squared test on discretized features

Unsupervised Model-Based Baselines:

- D3 (Discriminative Drift Detector) Gözüağık et al. (2019)
- MD3 (Margin Density Drift Detection) Sethi and Kantardzic (2017)
- Margin density-based detectors Ditzler and Polikar (2011)

3.6 Evaluation Metrics

According to Lu et al. (2018) and Souza et al. (2020a), **Drift Detection Performance** metrics include:

- Detection Delay: Time between true drift occurrence and detection
- False Positive Rate (FPR): Proportion of alarms when no drift exists
- True Positive Rate (TPR) / Sensitivity: Proportion of true drifts detected

- Precision: Proportion of alarms that are true drifts
- F1-Score: Harmonic mean of precision and recall
- AUC of Detection: Summary measure of detection performance

Model Performance Impact:

- Prequential Accuracy/F1: According to Gama et al. (2014), online accuracy using test-then-train protocol
- Recovery Time: Time to restore model performance after drift
- Mean Squared Error / RMSE: For regression tasks

3.7 Experimental Design

Experiment 1: Sensitivity to Drift Types

According to Lu et al. (2018), evaluate detection performance across abrupt, gradual, incremental, and recurring drift patterns on synthetic data with known drift points.

Experiment 2: Real-World Benchmark Comparison

According to Souza et al. (2020a), compare detection timeliness and model performance across real-world datasets against all baseline methods.

Experiment 3: Hyperparameter Sensitivity

Analyze impact of:

- Window sizes (N_{ref} , N_{cur})
- Classifier choice (logistic regression, random forest, XGBoost)
- Detection threshold (T_{drift})
- Monitoring frequency

Experiment 4: Scalability and Efficiency

Measure computational time and memory usage across varying data dimensionality and stream length.

Experiment 5: Adaptation Strategy Comparison

According to Gama et al. (2014), compare drift-triggered retraining vs. periodic retraining vs. no retraining in terms of downstream model performance.

3.8 Practical Considerations

Class Imbalance: When $|W_{ref}| \neq |W_{cur}|$, use:

- Stratified cross-validation
- Class-weighted classifiers
- According to Lopez-Paz and Oquab (2017), AUC-ROC rather than accuracy

Computational Efficiency:

- Use efficient classifiers (logistic regression, small random forests)
- Consider approximate methods for large windows
- According to Pan et al. (2020), implement incremental training where possible

Temporal Leakage Prevention:

- Use purged cross-validation to avoid temporal overlap between folds
- According to Souza et al. (2020a), ensure strict temporal ordering in train/test splits

Threshold Calibration:

- Use validation data with known drift to tune T_{drift}
- According to Pan et al. (2020), consider adaptive thresholds based on baseline variance

CHAPTER IV

RESULTS

CHAPTER V

DISCUSSION

CHAPTER VI

CONCLUSION

REFERENCES

- Baena-García, M., del Campo-Ávila, J., Fidalgo, R., Bifet, A., Gavaldà, R., and Morales-Bueno, R. (2006). Early drift detection method. In *Proceedings of the Fourth International Workshop on Knowledge Discovery from Data Streams (IWKDDs)*, pages 77–86.
- Bayram, F., Ahmed, B. S., and Kassler, A. (2022). From concept drift to model degradation: An overview on performance-aware drift detectors. *Knowledge-Based Systems*, 245:108632.
- Bifet, A. and Gavaldà, R. (2007). Learning from time-changing data with adaptive windowing. In *Proceedings of the 2007 SIAM International Conference on Data Mining (SDM)*, pages 443–448. SIAM.
- Dasu, T., Krishnan, S., Venkatasubramanian, S., and Yi, K. (2006). An information-theoretic approach to detecting changes in multi-dimensional data streams. In *Proceedings of the 38th Symposium on the Interface of Statistics, Computing Science, and Applications*.
- Ditzler, G. and Polikar, R. (2011). Hellinger distance based drift detection for nonstationary environments. In *IEEE Symposium on Computational Intelligence in Dynamic and Uncertain Environments (CIDUE)*, pages 41–48.
- Frias-Blanco, I., del Campo-Ávila, J., Ramos-Jiménez, G., Morales-Bueno, R., Ortiz-Díaz, A., and Caballero-Mota, Y. (2015). Online and non-parametric drift detection methods based on Hoeffding’s bounds. *IEEE Transactions on Knowledge and Data Engineering*, 27(3):810–823.
- Gama, J., Medas, P., Castillo, G., and Rodrigues, P. (2004). Learning with drift detection. In *Advances in Artificial Intelligence – SBIA 2004*, volume 3171 of *LNCS*, pages 286–295. Springer.

- Gama, J., Žliobaitė, I., Bifet, A., Pechenizkiy, M., and Bouchachia, A. (2014). A survey on concept drift adaptation. *ACM Computing Surveys*, 46(4):1–37.
- Gözüaık, Ö., Büyükakır, A., Bonab, H., and Can, F. (2019). Unsupervised concept drift detection with a discriminative classifier. In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management (CIKM)*, pages 2365–2368.
- Gözüaık, Ö. and Can, F. (2021). Concept learning using one-class classifiers for implicit drift detection in evolving data streams. *Artificial Intelligence Review*, 54(5):3725–3747.
- Gretton, A., Borgwardt, K. M., Rasch, M. J., Schölkopf, B., and Smola, A. (2012). A kernel two-sample test. *Journal of Machine Learning Research*, 13(25):723–773.
- Guo, C., Pleiss, G., Sun, Y., and Weinberger, K. Q. (2017). On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, volume 70 of *PMLR*, pages 1321–1330.
- Harries, M. (1999). Splice-2 comparative evaluation: Electricity pricing. Technical report, University of New South Wales.
- Hinder, F., Vaquet, V., Brinkrolf, J., and Hammer, B. (2024). One or two things we know about concept drift—a survey on monitoring in evolving environments. *Frontiers in Artificial Intelligence*, 7:1330257.
- Hsieh, C.-H., Li, M.-C., and Lee, S.-J. (2021). Concept drift detection based on pre-clustering and statistical testing. *Journal of Internet Technology*, 22(2):287–296.
- Krawczyk, B., Minku, L. L., Gama, J., Stefanowski, J., and Woźniak, M. (2017). Ensemble learning for data stream analysis: A survey. *Information Fusion*, 37:132–156.
- Kuncheva, L. I. (2013). Change detection in streaming multivariate data using likelihood detectors. *IEEE Transactions on Knowledge and Data Engineering*, 25(5):1175–1180.
- Liu, A., Song, Y., Zhang, G., and Lu, J. (2018). Accumulating regional density dissimilarity for concept drift detection in data streams. *Pattern Recognition*, 76:256–272.

- Liu, F., Xu, W., Lu, J., Zhang, G., Gretton, A., and Sutherland, D. J. (2020). Learning deep kernels for non-parametric two-sample tests. In *International Conference on Machine Learning (ICML)*, volume 119 of *PMLR*, pages 6316–6326.
- Long, M., Zhu, H., Wang, J., and Jordan, M. I. (2016). Unsupervised domain adaptation with residual transfer networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 29.
- Lopez-Paz, D. and Oquab, M. (2017). Revisiting classifier two-sample tests. In *International Conference on Learning Representations (ICLR)*.
- Lu, J., Liu, A., Dong, F., Gu, F., Gama, J., and Zhang, G. (2018). Learning under concept drift: A review. *IEEE Transactions on Knowledge and Data Engineering*, 31(12):2346–2363.
- Lukats, D., Zielinski, O., Hahn, A., and Stahl, F. T. (2025). A benchmark and survey of fully unsupervised concept drift detectors on real-world data streams. *International Journal of Data Science and Analytics*, 19:1–31.
- Lundberg, S. M. and Lee, S.-I. (2017). A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems 30 (NIPS)*, pages 4765–4774.
- Ovadia, Y., Fertig, E., Ren, J., Nado, Z., Sculley, D., Nowozin, S., Dillon, J., Lakshminarayanan, B., and Snoek, J. (2019). Can you trust your model’s uncertainty? evaluating predictive uncertainty under dataset shift. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, pages 13991–14002.
- Page, E. S. (1954). Continuous inspection schemes. *Biometrika*, 41(1/2):100–115.
- Pan, J., Pham, V., Doroudian, M., and Arora, N. (2020). MaLTA: Machine learning for user targeting at uber. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 3475–3483.
- Qian, H. and Rao, S. (2021). Managing dataset shift by adversarial validation for credit scoring. *arXiv preprint arXiv:2112.04471*.

- Raab, C., Heusinger, M., and Schleif, F.-M. (2020). Reactive soft prototype computing for concept drift streams. *Neurocomputing*, 416:340–351.
- Rabanser, S., Günnemann, S., and Lipton, Z. C. (2019). Failing loudly: An empirical study of methods for detecting dataset shift. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, pages 1396–1408.
- Renza, D. and Moya-Albor, E. (2024). Adversarial validation approach to concept drift detection in automated machine learning. *Applied Sciences*, 14(3):1178.
- Sculley, D., Holt, G., Golovin, D., Davydov, E., Phillips, T., Ebner, D., Chaudhary, V., Young, M., Crespo, J.-F., and Dennison, D. (2015). Hidden technical debt in machine learning systems. In *Advances in Neural Information Processing Systems 28 (NIPS)*, pages 2503–2511.
- Senoner, J., Netland, T. H., and Korherr, B. (2023). Using machine learning to predict and monitor operational performance in order picking. *International Journal of Production Research*, 61(23):8207–8227.
- Sethi, T. S. and Kantardzic, M. (2017). On the reliable detection of concept drift from streaming unlabeled data. *Expert Systems with Applications*, 82:77–99.
- Shang, D., Zhang, G., and Lu, J. (2020). Fast concept drift detection using unlabeled data. In *Developments of Artificial Intelligence Technologies in Computation and Robotics*, pages 1053–1060. World Scientific.
- Shimodaira, H. (2000). Improving predictive inference under covariate shift by weighting the log-likelihood function. *Journal of Statistical Planning and Inference*, 90(2):227–244.
- Souza, V. M. A., Reis, D. M., Maletzke, A. G., and Batista, G. E. A. P. A. (2020a). Challenges in benchmarking stream learning algorithms with real-world data. *Data Mining and Knowledge Discovery*, 34(6):1805–1858.

- Souza, V. M. A., Reis, D. M., Maletzke, A. G., and Batista, G. E. A. P. A. (2020b). Efficient unsupervised drift detector for fast and high-dimensional data streams. *Knowledge and Information Systems*, 62(4):1357–1385.
- Street, W. N. and Kim, Y. (2001). A streaming ensemble algorithm (SEA) for large-scale classification. In *Proceedings of the Seventh ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 377–382.
- Tsybmal, A. (2004). The problem of concept drift: Definitions and related work. Technical Report TCD-CS-2004-15, Trinity College Dublin, Department of Computer Science.
- Vergara, A., Vembu, S., Ayhan, T., Ryan, M. A., Homer, M. L., and Huerta, R. (2012). Chemical gas sensor drift compensation using classifier ensembles. *Sensors and Actuators B: Chemical*, 166:320–329.
- Vitter, J. S. (1985). Random sampling with a reservoir. *ACM Transactions on Mathematical Software*, 11(1):37–57.
- Webb, G. I., Hyde, R., Cao, H., Nguyen, H. L., and Petitjean, F. (2016). Characterizing concept drift. *Data Mining and Knowledge Discovery*, 30(4):964–994.
- Widmer, G. and Kubat, M. (1996). Learning in the presence of concept drift and hidden contexts. *Machine Learning*, 23(1):69–101.
- Žliobaitė, I. (2010). Learning under concept drift: An overview. *arXiv preprint arXiv:1010.4784*.

BIOGRAPHY

NAME	Ms. Nusrat Begum
DATE OF BIRTH	21 February 1998
PLACE OF BIRTH	Chattogram, Bangladesh
INSTITUTIONS ATTENDED	Stamford International University, 2017–2022 Bachelor of Business Administration (International Bus Mahidol University, 2024–2026 Master of Science (Computer Science) Mahidol University
E-MAIL	nusrat.beg@student.mahidol.edu